

Model Logs

Dataset:

<https://www.kaggle.com/datasets/xainano/handwrittenmathsymbols/code>

Approach:

1. Read an image of a linear equation.
2. Get the bounding box for each individual characters.
3. Extract the bounding box from the image.
4. Preprocess every extracted sub-image – S'
5. Load the saved model which is stored in pickle format.
6. Run every S' through model.
7. Get the predictions for each image.
8. Use the sympy for solving the equation.


Implementation:

- Fully connect Deep neural network
- RELU activation function
- Softmax activation function
- Cross entropy loss function
- Convolutional neural network
- Flattening layer
- MaxPool layer.
- Sequential Model

Challenges:

- The biggest hurdle was implementing the convolution neural network especially the backpropagation algorithm for convolution algorithm.
- Even though the algorithm is straightforward it was tough to implement as we were frequently encountering issue in managing the gradients for each activation layer.
- Vectorizing the code was another challenge to solve, NumPy really helped in performing the vectorization and speeding up the code.
- In order to validate the correct working of each layer, we tried training each layer individually on the known dataset such as MNIST, after we validated that the layer is able to learn features by looking into it's validation accuracy during training we integrated it our sequential model. We frequently faced issues while integrating the layers together too.

Data Distribution:



The class distribution is not balanced, so we may need to use data augmentation to balance the dataset.

Main classes include:

Digits: 0 to 9

Characters: - a - z (some include upper case characters)

Math operators: - +, -, /, *, =

It also includes basic Greek alphabet symbols like: alpha, beta, gamma, mu, sigma, phi and theta.

English alphanumeric symbols are included.

All math operators, set operators.


Basic pre-defined math functions like: log, lim, cos, sin, tan.

Number of classes - 82

Features - As it is a gray scale image dataset the features are just the number of pixels present in an image, so it will be 2025 pixels / features.

Data Preprocessing:

We didn't choose all the 82 classes from the original dataset; we chose the below 20 classes for our use case.



Results & Analysis:

We've implemented 4 models with different architecture & parameters:

Model 1: Flatten1, Dense1, Softmax, Cross Entropy

Model Info

Layer Name-> Flatten 1

Layer Name -> Dense_1



Weights shape -> (2025, 20)

Model	Epochs	Batch Size	Learning Rate	Train Accuracy (%)	Validation Accuracy (%)	Train Loss	Validation Loss	Test Accuracy (%)	Test Loss (%)
Model 1	5	16	0.0001	4.86	5.04	28.75	28.43	-	-
Model 1	5	16	0.001	7.07	8.00	24.58	24.02	-	-
Model 1	20	16	0.001	15.70	15.99	18.19	18.38	-	-
Model 1	20	16	0.001	18.79	18.78	16.23	16.21	-	-
Model 1	30	16	0.01	55.83	56.54	5.51	5.38	-	-
Model 1	60	16	0.01	66.40	65.95	3.67	3.79	-	-
Model 1	100	16	0.01	72.45	71.55	2.70	2.89	-	-
Model 1.1	150	16	0.01	76.34	74.89	2.14	2.36	75.07	2.34
Model 1.2	20	16	0.01	84.88	85.01	0.60	0.61	84.02	0.62



Class	Precision	Recall	F1-Score	Support
0	0.92	0.95	0.93	1368
1	0.83	0.91	0.87	1385
2	0.80	0.80	0.80	1452
3	0.90	0.90	0.90	1439
4	0.76	0.83	0.79	1391
5	0.89	0.80	0.84	702



6	0.92	0.88	0.90	646
7	0.89	0.84	0.87	576
8	0.84	0.79	0.82	637
9	0.83	0.85	0.84	764
div	0.99	0.44	0.60	179
rightarrow	0.91	0.80	0.85	352
times	0.87	0.81	0.84	702
u	0.62	0.42	0.50	256
v	0.79	0.69	0.74	295
w	0.72	0.64	0.68	104
y	0.78	0.84	0.81	1329
z	0.76	0.72	0.74	1135
+	0.80	0.83	0.82	1378
-	0.94	0.99	0.96	1385
Accuracy	0.84			


Graphs for model 1.1:



Graphs for Model 1.2 (Single layer DNN with :







Observations for model (Accuracy 84%):

1. Weight initialization: Without glorot/xavier weight initialization the model took more number of epochs for convergence compared to glorot/Xavier initialization.
2. As it is a single layer NN, it is the fastest compared to all models in training time.
3. The Model 1.2 was stuck in local minima i.e it's accuracy didn't improve above 74% after 100 epochs.
4. But the same model (Model 1.2) with glorot init converged faster at 20 epochs, with validation and testing accuracy of 84% at the end of 20 epoch.

Model 2: Dense1, ReLu, Dense2, Softmax, Cross Entropy

Layer information

Number of Layers 4

Layer 1

Layer Name -> flatten_1

Layer 2

Layer Name -> Dense_1

Weights shape -> (2025, 50)

Layer 3

Layer Name -> relu_1

Layer 4


Layer Name -> Dense_2

Weights shape -> (50, 20)

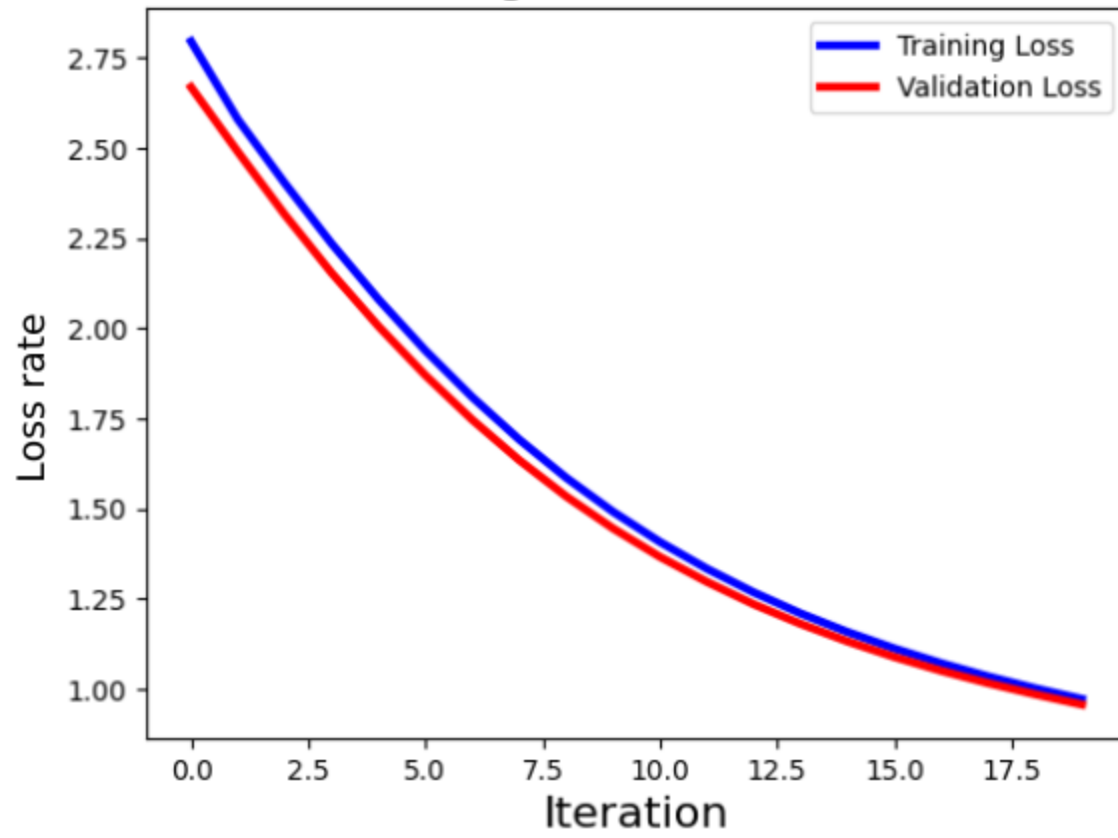
Model	Epochs	Batch Size	Learning Rate	Train Accuracy (%)	Validation Accuracy (%)	Train Loss	Validation Loss	Test Accuracy (%)	Test Loss (%)
Model 2	5	16	0.0001	4.86	5.04	28.75	28.43	-	-
Model 2.1	5	16	0.001	7.07	8.00	24.58	24.02	-	-
Model 2.1	20	16	0.001	15.70	15.99	18.19	18.38	-	-
Model 2.2	20	16	0.001	18.79	18.79	16.23	16.21	-	-

Class	Precision	Recall	F1-Score	Support
0	0.83	0.92	0.87	1368
1	0.73	0.85	0.79	1385
2	0.70	0.72	0.71	1452
3	0.79	0.85	0.82	1439
4	0.64	0.73	0.68	1391
5	0.83	0.69	0.75	702
6	0.79	0.74	0.77	646
7	0.83	0.64	0.72	576
8	0.76	0.61	0.67	637
9	0.72	0.67	0.69	764
div	0.83	0.03	0.05	179
rightarrow	0.77	0.45	0.57	352
times	0.78	0.66	0.72	702
u	0.42	0.23	0.30	256
v	0.75	0.30	0.43	295
w	0.88	0.21	0.34	104


y	0.63	0.76	0.69	1329
z	0.64	0.62	0.63	1135
+	0.67	0.74	0.70	1378
-	0.88	0.97	0.92	1385
Accuracy	0.73			
Macro Avg	0.74	0.62	0.64	17475
Weighted Avg	0.74	0.73	0.72	17475




Training vs Validation loss



Learning Curve





Observations

1. Model is stuck at local minima (Model 2.1) but with gloriot init (Model 2.2) it improved.
2. Model 1.2 has better metrics than Model 2.2 which has the same hyperparameters but Model 2.2 has 2 Dense layers whereas Model 1.2 has only one dense layer

Model 3: Conv1, Relu1, Maxpool1, Flatten1, Dense1, Softmax, Cross Entropy

Layer information

Layer 1

Layer Name -> conv_1

Number of filters -> 8

Filters shape -> (8, 1, 3, 3)

Stride -> 1

Layer 2

Layer Name -> relu_1

Layer 3

Layer Name -> max_pool_1

Pool size -> 2

Stride -> 2

Layer 4

Layer Name -> flatten_1

Layer 5



Layer Name -> Dense_1



Weights shape -> (3528, 20)


Model	Epochs	Batch Size	Learning Rate	Train Accuracy (%)	Validation Accuracy (%)	Train Loss	Validation Loss	Test Accuracy (%)	Test Loss (%)
Model 3	11	64	0.01	66.42	67.18	1.32	1.28	68.00	-
Model 3	20	64	0.01	81.17	81.40	0.73	0.73	81.00	-

Class	Precision	Recall	F1-Score	Support
0	0.89	0.93	0.91	1368
1	0.78	0.86	0.82	1385
2	0.76	0.76	0.76	1452
3	0.88	0.87	0.87	1439

4	0.73	0.82	0.77	1391
5	0.87	0.78	0.82	702
6	0.86	0.87	0.87	646
7	0.83	0.78	0.81	576
8	0.80	0.72	0.76	637
9	0.78	0.81	0.79	764
div	0.99	0.39	0.55	179
rightarrow	0.86	0.76	0.81	352
times	0.85	0.77	0.81	702
Accuracy:	0.81			







Observations for Model 3:

1. Took more time to train as convolution operation was involved.
2. Achieved similar metric to Model 1.2 but took longer time for each epoch.
3. Observed a spike in loss, during the 11 epoch.

Model 4: Conv2, Relu1, MaxPool1, Conv1, Relu2, MaxPool2, Flatten1, Dense1, Softmax, Cross Entropy

Layer information

Number of Layers 8

Layer 1

Layer Name -> conv_1

Number of filters -> 8

Filters shape -> (8, 1, 3, 3)

Stride -> 1

Layer 2

Layer Name -> relu_1

Layer 3

Layer Name -> max_pool_1

Pool size -> 2

Stride -> 2

Layer 4

Layer Name -> conv_2

Number of filters -> 12

Filters shape -> (12, 8, 5, 5)

Stride -> 1

Layer 5

Layer Name -> relu_2

Layer 6

Layer Name -> max_pool_2

Pool size -> 2

Stride -> 2

Layer 7

Layer Name -> flatten_1

Layer 8


Layer Name -> Dense_1



Weights shape -> (768, 20)



Model	Epochs	Batch Size	Learning Rate	Train Accuracy (%)	Validation Accuracy (%)	Train Loss	Validation Loss	Test Accuracy (%)	Test Loss (%)
Model 4	20	128	0.1	83.92	83.99	0.5917	0.5874	83.45	0.5988

Class	Precision	Recall	F1-Score	Support
0	0.89	0.93	0.91	1368
1	0.83	0.90	0.86	1385
2	0.81	0.79	0.80	1452
3	0.91	0.89	0.90	1439
4	0.79	0.83	0.81	1391
5	0.85	0.80	0.82	702
6	0.86	0.86	0.86	646
7	0.76	0.71	0.73	576
8	0.82	0.75	0.78	637
9	0.73	0.85	0.79	764
div	0.94	0.60	0.73	179
rightarrow	0.86	0.78	0.82	352
times	0.88	0.85	0.87	702
u	0.70	0.54	0.61	256
v	0.74	0.72	0.73	295

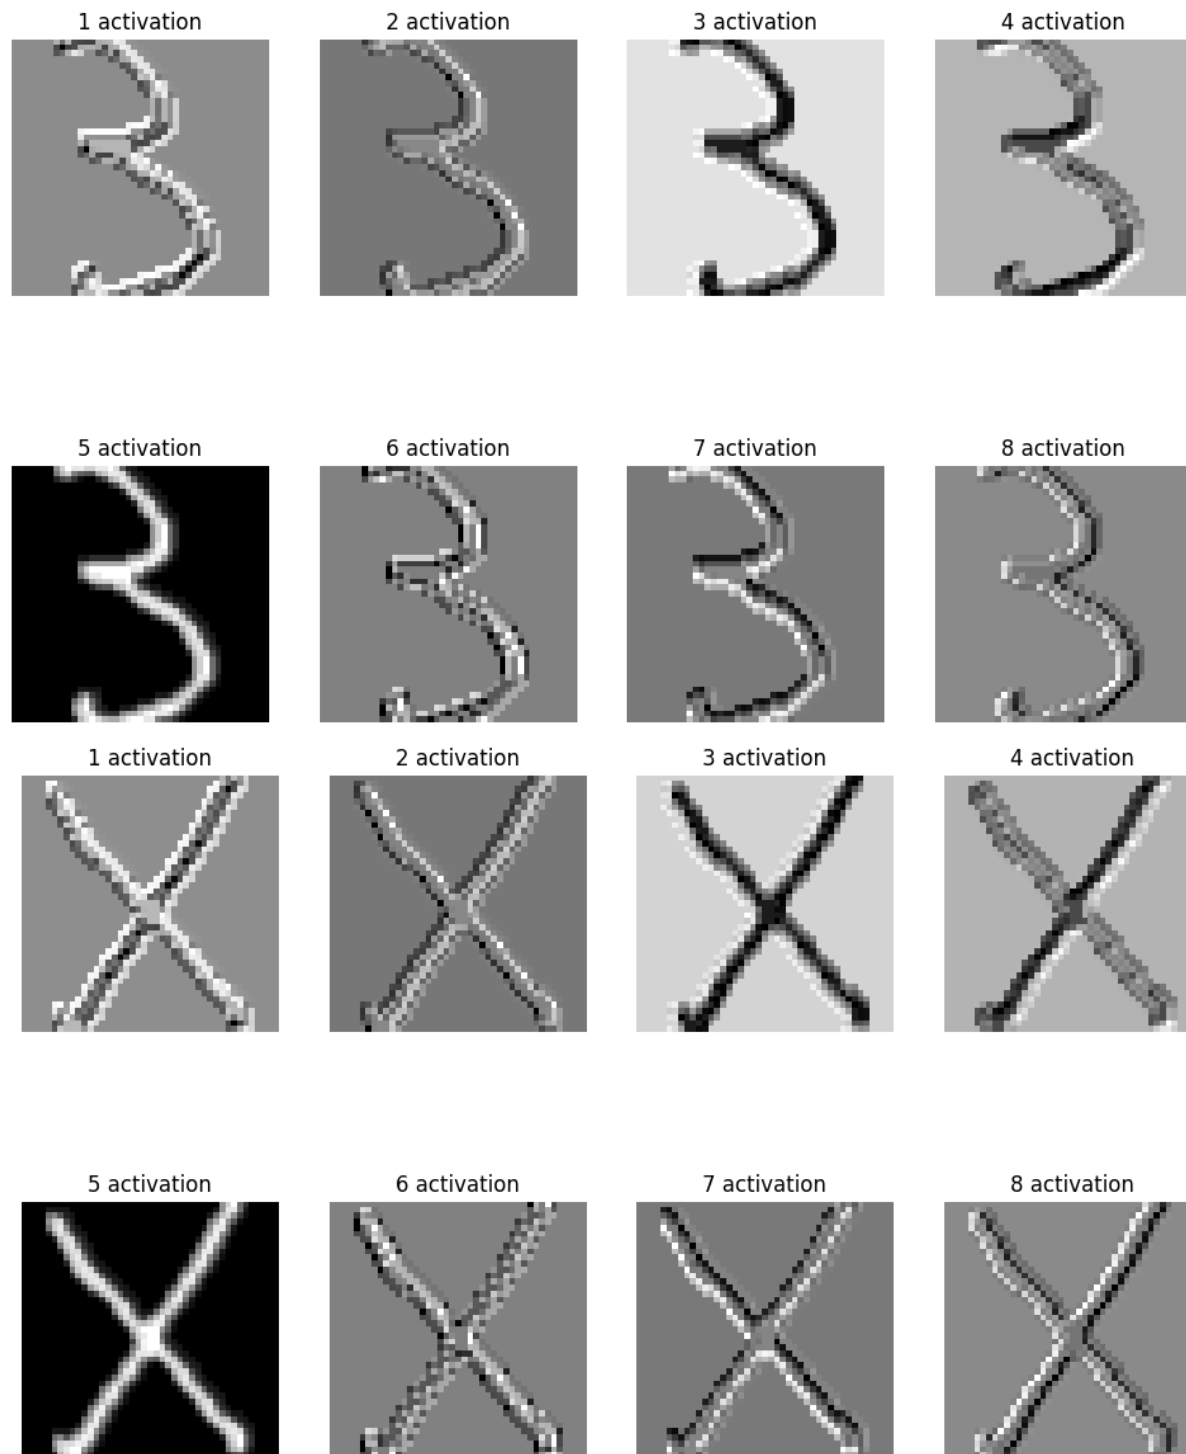
w	0.76	0.66	0.71	104
y	0.78	0.78	0.78	1329
z	0.74	0.72	0.73	1135
+	0.87	0.88	0.88	1378
-	0.95	0.99	0.97	1385
Accuracy:			0.83	17475
Macro Avg:	0.82	0.79	0.8	17475
Weighted Avg:	0.83	0.83	0.83	17475









Convolutional Layer 1 Output Activations:



1 activation



2 activation




3 activation




4 activation




5 activation




6 activation







7 activation

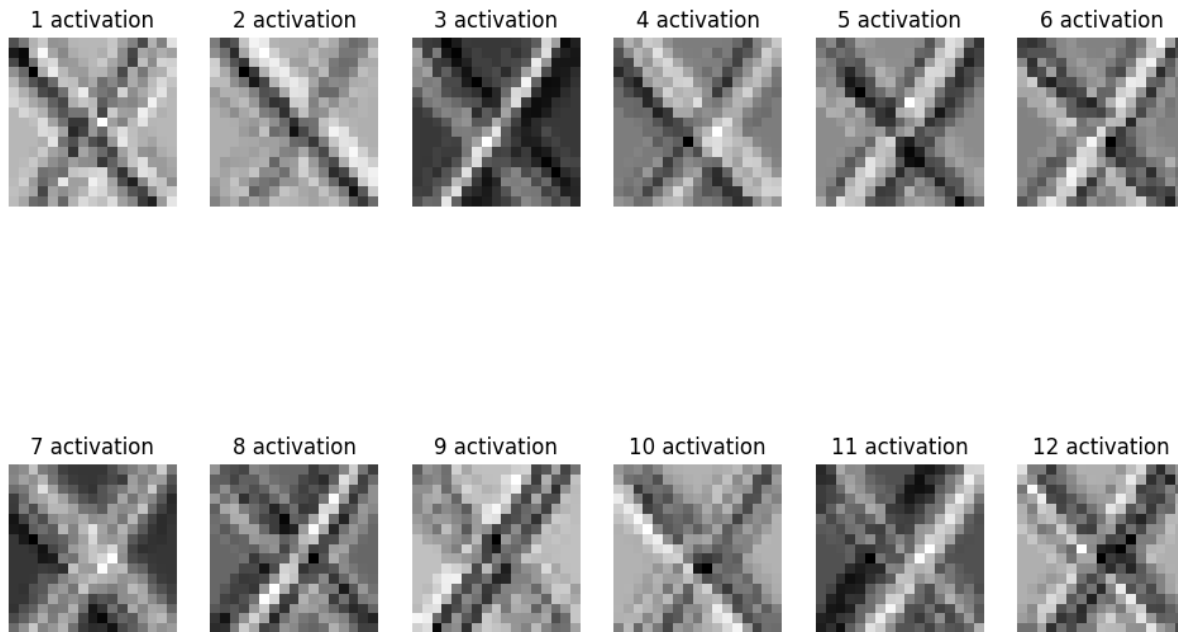


8 activation



Convolutional Layer 2 Output Activations:

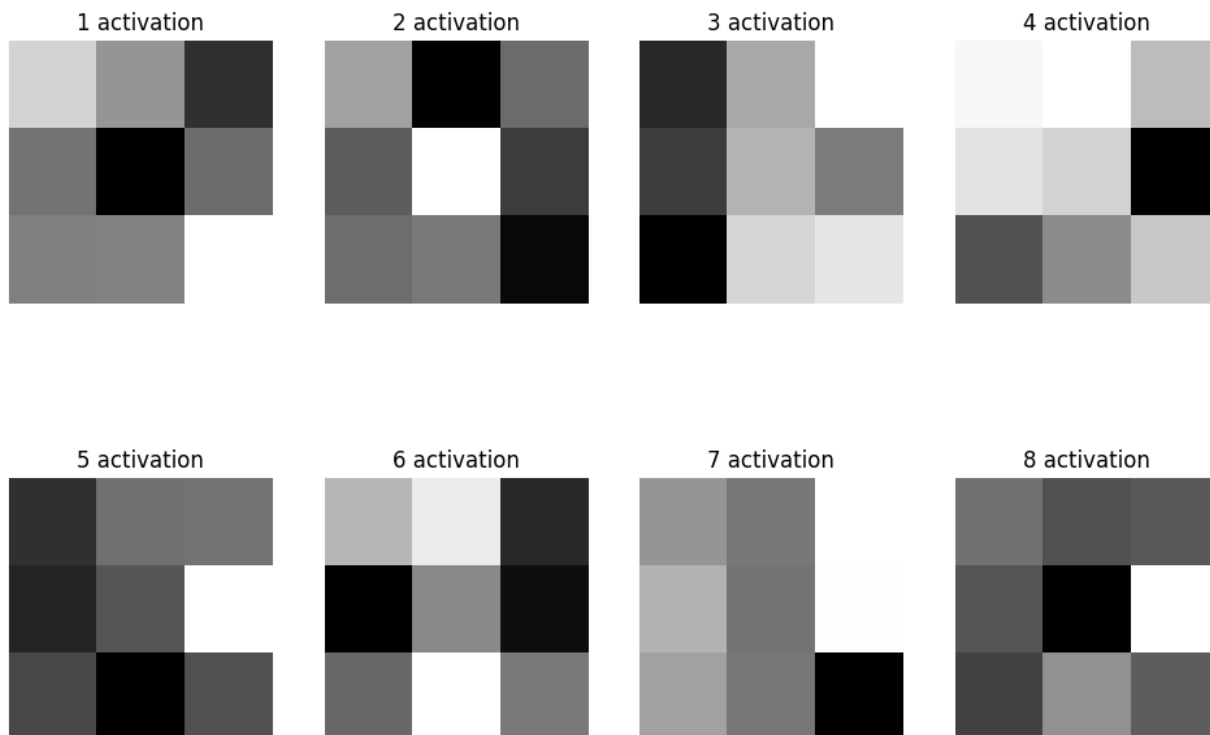




Filter Activations

Convolutional Layer 1 Filter Output



8 Filters of Size 1:3:3 (Channels, Width, Height)



Convolutional Layer 2 Filter Output

12 Filters of size 8:5:5 (Channels, Width, Height)

Only visualizing 12:1:5:5



Conclusion and future work:

1. The best performing model on real world images didn't work as well as it did on test set. So there might be some issue while preprocessing real-images before feeding into the model. So, the model is overfitting as it's not working on real world images.
2. More work need to be done on dataset preparation as the model on real world images weren't good, so we need to augment the images for the classes which are of less numbers.
3. Forward propagation is vectorized in CNNs but during backward propagation we can vectorize the code even for calculating the gradients for each channel.
4. Currently we only solve linear equations with single variables, we can work on extending the equation solver for quadratic and trigonometric equations.
5. CNN layers are pretty slow, so we can look into improving the training speed of these layers by using Python and vectorizing the code over the image channels.

References:

1. <https://www.parasdahal.com/softmax-crossentropy>
2. https://jmlb.github.io/ml/2017/12/26/Calculate_Gradient_Softmax/
3. https://www.youtube.com/watch?v=5-rVLSc2XdE&ab_channel=SmartAlphaAI
4. <https://pyimagesearch.com/2021/05/06/understanding-weight-initialization-for-neural-networks/>
5. <https://www.youtube.com/watch?v=3TdBtI9dh2I>
6. https://deeplearning.cs.cmu.edu/F21/document/recitation/Recitation5/CNN_Backprop_Recitation_5_F21.pdf
7. <https://www.educative.io/answers/how-to-backpropagate-through-max-pooling-layers>
8. <https://pyimagesearch.com/2015/04/20/sorting-contours-using-python-and-opencv/>